Revisiting mortality prediction in ICU patients with diabetic neuropathy: a 24-hour landmark study using MIMIC-IV

BMI5101 Group Project (P16)

[Author names and affiliations]

October 22, 2026

Abstract

Background. Huang et al. (2025) built machine-learning models for mortality in intensive care unit (ICU) patients with diabetic neuropathy (DN) from MIMIC-IV. After LASSO selection of 12 of 55 candidate predictors, random forest (RF) had the highest reported validation AUC (0.780) among RF, XGBoost, support vector machine and logistic regression (LR). However, inclusion used information from after the first ICU day, the outcome was not consistently defined, and models were chosen on one 263-patient validation set after variable selection on the whole training set, without nested validation.

Methods. Keeping the benchmark's framework, we used MIMIC-IV v3.1 to build a landmark model predicting at 24 h after ICU admission from first-day measurements and retrospectively coded comorbidities. We retained patients without a white-cell differential, separated chronic from acute comorbidity codes and defined death within 365 days. Models were evaluated by 5×10 repeated nested cross-validation with imputation, LASSO and tuning inside each outer training fold, and chosen by a pre-specified rule; the selected model was frozen and evaluated once on a held-out set. A reference cohort with the benchmark's criteria tested the original pipeline. As exploratory extensions, first-day trajectory features, prior-diagnosis embeddings and stacking were compared on the development set.

Results. In the reference cohort, the benchmark's RF reproduced its training AUC of 0.999 but reached only 0.724–0.767 under nested cross-validation, against 0.791 for the best model. In the landmark cohort (2,557 patients, 26.6% mortality), six configurations met the pre-specified criteria for selection and L2-penalised LR was selected (nested-CV AUROC 0.801). Its nested-CV AUPRC was 0.594 (prevalence 0.266). On the 512-patient held-out set its AUROC was 0.750 (95% CI 0.702–0.797), its AUPRC 0.533 (0.459–0.611) and its calibration slope 0.82. Red cell distribution width, Charlson comorbidity index and age were the leading predictors. Stacking changed AUROC by 0.005 over LR (95% CI -0.001 to 0.011), an uncertain gain, and trajectory features and diagnosis embeddings gave no gain.

Conclusions. Within the benchmark's LASSO-based framework, we found no evidence that RF outperforms LR. The landmark model has a defined time of use; the real-time availability of its inputs, its calibration in other settings and its external validity remain to be established.

Keywords: diabetic neuropathy; ICU; mortality prediction; MIMIC-IV; landmark model; nested cross-validation; LASSO; calibration

1 Introduction

1.1 Background

Diabetic neuropathy is the most common chronic complication of diabetes and is associated with foot ulceration, amputation, cardiovascular autonomic dysfunction and increased mortality (Feldman et al., 2019; Pop-Busui et al., 2017). Patients with DN who are admitted to the ICU typically carry a heavy burden of long-standing diabetes and comorbid cardiovascular and renal disease. Identifying those at highest risk of death in the year after ICU admission could support goals-of-care discussions and targeted follow-up.

1.2 The benchmark study

Huang et al. (2025) used the MIMIC-IV database (Johnson et al., 2023) to predict death in ICU patients with DN. Their framework had five steps: an ICD-defined DN cohort; 55 candidate predictors from the first ICU day (demographics, comorbidities, laboratory values, vital signs, severity scores and inflammatory ratios); LASSO selection of 12 predictors (Tibshirani, 1996); RF, XGBoost, SVM and LR models; and SHAP interpretation of the best model (Lundberg & Lee, 2017), supported by decision-curve and clinical-impact analyses. RF was reported as the best model (validation AUC 0.780), and RDW, neutrophil count, the Charlson comorbidity index (CCI), chloride and prothrombin time were the most important predictors. The framework is reasonable and widely used, but several design choices limit what its results can show (Collins et al., 2015).

1.3 Issues and our modifications

Table 1 lists seven issues in the benchmark, what each implies, and how we changed the analysis. The issues differ in the strength of evidence: I1, I2, I3, I6 and the single validation set in I5 are stated in the benchmark; the missing outcome window, hypertension definition, calibration and tuning grids are not reported; and the mixed differential units and the positive class of the threshold metrics (I4, I7) are our inferences from its tables, which its code would be needed to confirm. The rest of the paper follows this table: Section 2 implements each modification, Section 3 gives the evidence, and Section 4 summarises the effect of each change.

Table 1: Issues in the benchmark and our modifications.

#	Issue in the benchmark	Consequence	Our modification	Sections
I1	Inclusion required a hospital stay longer than 48 h.	Length of stay is unknown on the first ICU day; early deaths and early discharges are excluded, and the model has no defined time of use.	Fixed prediction time at 24 h after ICU admission (landmark); all adults alive and in the ICU at that time are included.	2.2, 3.1
I2	All eight complete-blood-count items were required.	About one third of eligible patients are excluded, almost all because no white-cell differential was performed.	These patients are kept; a missing-differential indicator is added and values are imputed within folds.	2.3, 3.1, 3.6

#	Issue in the benchmark	Consequence	Our modification	Sections
I3	Comorbidities were taken from all codes of the index admission.	Untimed codes for acute events of the current stay are counted as pre-existing disease.	Comorbidities from prior admissions plus chronic index-admission codes; CCI without its age component.	2.3, 3.6
I4	The outcome was described inconsistently (in-hospital death in the abstract, a non-hospital death within one year in a SHAP example) and was not defined in the Methods; white-cell differential units appear mixed; the hypertension definition was not given.	Mortality cannot be interpreted; differential-based predictors may reflect test format; a narrow hypertension definition creates a spurious protective effect.	Death within 365 days of ICU admission; absolute cell counts; full Quan–Elixhauser hypertension definition.	2.3, 3.2
I5	Models were compared on one 263-patient validation set; LASSO was run once on the whole training set, and 10-fold cross-validation for tuning was not nested within an outer validation loop.	Model differences are smaller than the uncertainty of a single small validation set, and performance estimates are optimistic (Cawley & Talbot, 2010; Varma & Simon, 2006).	Repeated nested cross-validation with every data-dependent step inside the folds; corrected resampled t-test; pre-specified selection rule; frozen model evaluated once on a held-out set.	2.5, 3.3, 3.4
I6	RF was chosen despite a training AUC of 0.999 against a validation AUC of 0.780, a gap the authors themselves reported as overfitting; tuning grids were not reported.	The selected model is likely overfitted, and comparator models may not have been tuned comparably.	Every model family tuned on stated grids inside the folds and compared side by side with the benchmark's parameters.	2.4, 3.3
I7	Performance was reported as AUC, threshold metrics, decision curves and clinical impact curves on the validation set, without calibration or Brier score; the threshold metrics appear consistent only with survival as the positive class.	Whether predicted risks are accurate is unknown, and the classification results do not describe the detection of deaths.	Calibration intercept and slope, Brier and scaled Brier scores and AUPRC added; death as the positive class with a threshold from out-of-fold predictions; decision curves on cross-validated and frozen held-out predictions.	2.6, 3.5

1.4 Contributions

This study makes four contributions:

1. **A re-evaluation of the benchmark framework.** Repeated nested cross-validation, a corrected test and a pre-specified selection rule are applied to the benchmark's own model configurations and to tuned alternatives.
2. **A landmark model with a defined time of use.** The model predicts at 24 h after ICU admission, conditional on survival in the ICU to that time, retains patients without a white-cell differential and separates prior-admission and chronic index-admission diagnoses from acute ones; it was frozen before a single confirmatory held-out evaluation.
3. **An empirical check of the benchmark's issues.** Its overfitting was reported by its authors; our contribution is to quantify and explain it. For each issue in Table 1 we report how much the modification changed the cohort, the predictors or the performance, including the modifications that made little difference.

4. **An exploratory test of proposed extensions.** First-day trajectory features, prior-diagnosis representations and a stacking ensemble are compared with matched controls on the same folds.

2 Methods

2.1 Data source and the framework we build on

We used MIMIC-IV version 3.1 (Johnson et al., 2023), which contains de-identified records of patients admitted to the Beth Israel Deaconess Medical Center (Boston, USA) between 2008 and 2022; the benchmark used an earlier v2.x release covering 2008–2019. Severity scores, the comorbidity index and first-day laboratory and vital-sign summaries were computed with the official MIT-LCP `mimic-code` concepts. Analyses were run in Python 3.11 with DuckDB 1.5.6, scikit-learn 1.9, xgboost 3.4 and shap 0.51. In line with the PhysioNet data use agreement, only aggregate results are reported.

We kept the benchmark's framework: the ICD-defined DN population, the 55 candidate predictors from the first ICU day, LASSO selection at λ.1se, the four model families (RF, XGBoost, SVM, LR), the benchmark's reported hyperparameters as one model configuration, SHAP interpretation and decision-curve analysis. The changes are those listed in Table 1 and are described in Sections 2.2–2.6.

2.2 Landmark cohort (I1)

The study population consisted of patients with at least one hospital admission carrying an ICD-9 or ICD-10 code for DN and an ICU stay. For each patient, the **index admission** was the first admission with both a DN code and an ICU stay, and the index ICU stay was the first ICU stay of that admission. The benchmark's wording ("first ICU admission at first hospitalization") admits several readings; readings that require the very first admission or the first-ever ICU stay to qualify exclude eligible patients because of their care history, and are summarised in Supplementary Table S1.

Prediction was made at a fixed **landmark of 24 h after ICU admission** (van Houwelingen, 2007). Patients aged 18 years or older who were still in the ICU and alive at that time were included. Patients whose first ICU stay ended within 24 h were excluded, and a later ICU stay of the same admission was not used in its place. No length-of-stay or blood-count criterion was applied, because neither is known, or clinically required, at the time of prediction. The landmark cohort was split once into a stratified 80% training set and a 20% held-out set.

For comparison, a **reference cohort** applied the benchmark's criteria (hospital stay longer than 48 h, age over 18 years and all eight complete-blood-count items) to the same index admissions; it is analysed in Section 3.3 and the Supplement.

2.3 Outcome and predictors (I2, I3, I4)

Outcome (I4). The outcome was all-cause death on a date between the day of ICU admission and 365 days later, using the in-hospital time of death when available and the date of death from state records otherwise, compared at date granularity. MIMIC-IV records out-of-hospital deaths up to one year after a patient's last discharge, so deaths within the 365-day window are captured to the extent recorded in hospital and state death records; the window boundary is resolved only to the day. The model therefore

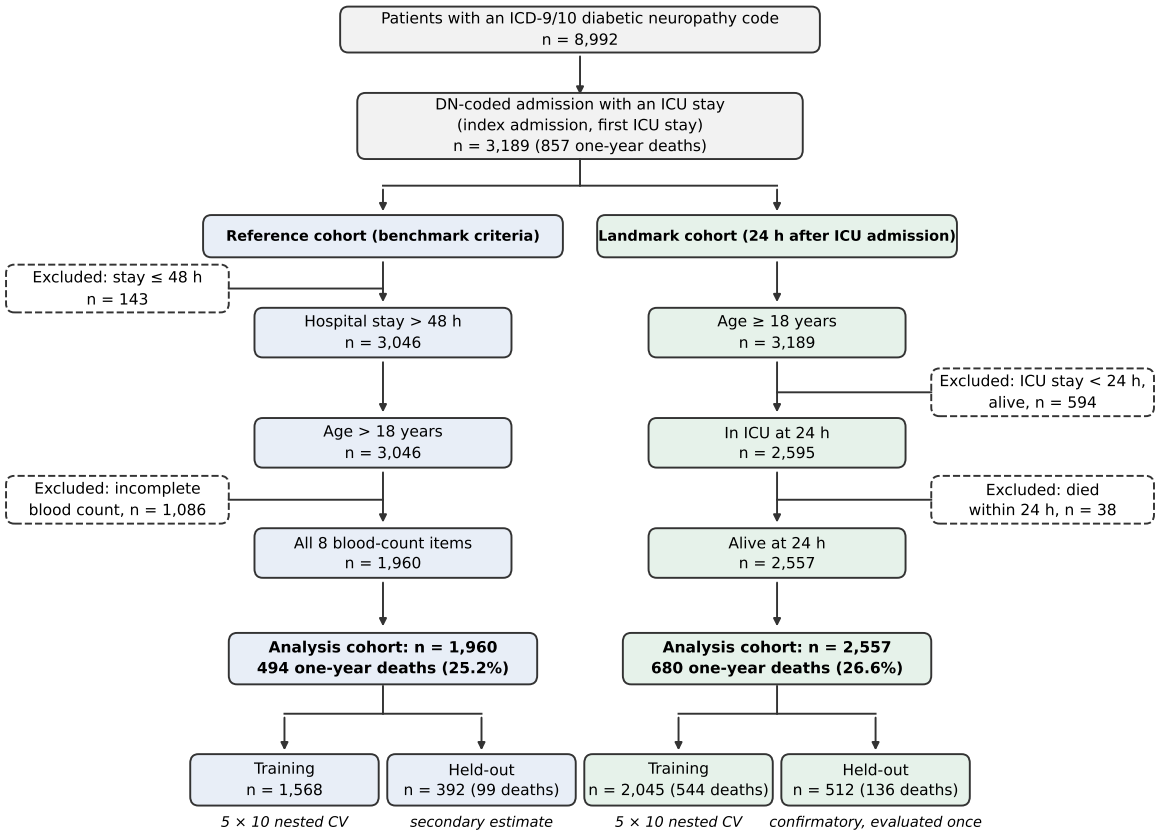


Figure 1: Cohort flow. Both cohorts start from the same index admissions. The reference cohort applies the benchmark's inclusion criteria; the landmark cohort includes all adults alive and in the ICU 24 h after admission.

estimates, for patients alive and in the ICU at 24 h, the risk of death by day 365 after ICU admission; it does not apply to patients who die or leave the ICU within the first 24 h. Cause of death is not available.

Candidate predictors. We extracted the benchmark's 55 candidate predictors: age and sex; six comorbidities (myocardial infarction, congestive heart failure, chronic pulmonary disease, renal disease, hypertension, obesity); 31 laboratory summaries (glucose, anion gap, bicarbonate, creatinine, urea nitrogen, haemoglobin, platelets, white and red cells, RDW, absolute lymphocytes, neutrophils and monocytes, calcium, chloride, sodium, potassium, PTT, PT, INR); four vital signs (heart rate, mean blood pressure, respiratory rate, temperature); six severity scores (SAPS-II, APS-III, LODS, OASIS, GCS, CCI); and six inflammatory ratios (NLR, PLR, LMR, PNR, PWR, SII). Laboratory values were taken from 6 h before to 24 h after ICU admission and vital signs from the first 24 h. Windows were defined on the measurement (chart) time; we did not check whether each laboratory result had been reported by 24 h. Every laboratory item identifier was checked against the MIMIC-IV dictionaries.

Units and definitions (I4). White-cell differentials were taken as absolute counts (10^9/L) from the `mimic-code` blood-differential concept, so that no patient's value mixes absolute counts with percentages. Hypertension followed the Quan–Elixhauser definition including complicated hypertension (ICD-9 401–405; ICD-10 I10–I13, I15, I16) (Quan et al., 2005). Under ICD rules, hypertension with chronic kidney disease or heart failure is coded under the complicated codes alone, so an uncomplicated-only definition labels these patients as normotensive.

Comorbidity timing (I3). Diagnosis codes in MIMIC-IV carry no timestamps, so codes of the index admission may describe acute events of the current stay. In the landmark cohort, comorbidities were derived from all codes of prior admissions plus the index-admission codes that remained after removing acute-event codes (for example acute myocardial infarction, acute heart failure, acute stroke or transient ischaemic attack, aortic dissection, acute peptic ulcer and shock liver). The CCI was computed without its age component, because age enters the model separately, and the chronic-disease item of SAPS-II was recomputed over prior and index admissions.

White-cell differential (I2). Instead of excluding patients without a differential, we added one binary candidate predictor, `wbc_diff_missing`, equal to 1 when all three absolute cell counts were missing. The eight differential-derived variables (three cell counts and NLR, PLR, LMR, PNR, SII) were imputed within folds like every other variable. The landmark analysis therefore had 56 candidate predictors.

2.4 Models (I6)

Preprocessing. Within each training set, variables with more than 20% missing values were dropped (the eight differential-derived variables were exempt in the landmark cohort). Missing values were then imputed with the training mean for continuous variables and the training mode for binary variables and ordinal scores, and each variable was standardised to zero mean and unit variance using training statistics. LASSO, LR and SVM used standardised values; RF and XGBoost used imputed values on the original scale.

Logistic regression (final model). For patient i with predictor vector $\mathbf{x}_i$, the model is

$$p_i = \sigma(\beta_0 + \mathbf{x}_i^\top \boldsymbol{\beta}), \qquad \sigma(z) = \frac{1}{1+e^{-z}}.$$

LASSO selection. Variables were selected by L1-penalised LR (Tibshirani, 1996), which solves

$$\min_{\beta_0,\boldsymbol{\beta}} \; -\frac{1}{n}\sum_{i=1}^{n}\left[y_i \log p_i + (1-y_i)\log(1-p_i)\right] + \lambda\|\boldsymbol{\beta}\|_1,$$

with the intercept unpenalised. We used scikit-learn's `liblinear` solver, which minimises $\|\boldsymbol{\beta}\|_1 + C\sum_i \ell_i$, where ℓ_i is the log-loss of patient i. Dividing by nC shows that the two problems are identical when $C = 1/(n\lambda)$, and this conversion was used for every λ (implementation details in Supplementary S9).

The solver (newGLMNET) is a coordinate-descent method of the glmnet type (Friedman et al., 2010; Yuan et al., 2012). At each outer iteration it replaces the log-likelihood by a weighted least-squares approximation around the current estimate, with working responses z_i and weights $w_i = \tilde{p}_i(1-\tilde{p}_i)$, and then updates one coefficient at a time. The update for coefficient j is

$$\beta_j \leftarrow \frac{S\big(\frac{1}{n}\sum_i w_i x_{ij} r_{ij},\ \lambda\big)}{\frac{1}{n}\sum_i w_i x_{ij}^2}, \qquad S(z,\lambda) = \operatorname{sign}(z)\,(|z|-\lambda)_+,$$

where r_{ij} is the partial residual of z_i that excludes predictor j. The soft-threshold operator S returns exactly zero whenever $|z| \le \lambda$, that is, whenever the weighted association between predictor j and the partial residual is too weak to offset the penalty. Coefficients are therefore set exactly to zero rather than merely shrunk, and the predictors with non-zero coefficients form the selected set.

Choice of λ. Over a glmnet-style grid of λ values (Supplementary S9), stratified 10-fold cross-validation gave the mean held-out binomial deviance $\mathrm{CV}(\lambda)$ and its standard error $\mathrm{SE}(\lambda)$, the standard deviation across folds divided by $\sqrt{10}$. With $\lambda_{\min} = \arg\min_\lambda \mathrm{CV}(\lambda)$, the **$\lambda$.1se rule** takes the largest λ with $\mathrm{CV}(\lambda) \le \mathrm{CV}(\lambda_{\min}) + \mathrm{SE}(\lambda_{\min})$, the sparsest model whose deviance is within one standard error of the minimum. The predictors with non-zero coefficients at λ.1se, refitted on the whole training data, were retained.

L2-penalised and unpenalised LR. On the selected predictors, the L2-penalised LR solves

$$\min_{\beta_0,\boldsymbol{\beta}} \; \frac{1}{n}\sum_{i=1}^{n}\ell_i + \frac{\lambda_2}{2}\|\boldsymbol{\beta}\|_2^2.$$

scikit-learn minimises $\frac{1}{2}\|\boldsymbol{\beta}\|_2^2 + C\sum_i \ell_i$ (L-BFGS solver, intercept unpenalised), which is the same problem with $\lambda_2 = 1/(nC)$. C was tuned over $\{0.001, 0.01, 0.1, 1, 10, 100, \infty\}$, where $C = \infty$ removes the penalty. The **unpenalised LR**, which corresponds to the benchmark's LR fitted with R `glm`, was also fitted as a separate configuration.

Random forest. An RF averages B classification trees, each grown on a bootstrap sample with a random subset of predictors considered at each split (Breiman, 2001):

$$\hat{p}(\mathbf{x}) = \frac{1}{B}\sum_{b=1}^{B} T_b(\mathbf{x}),$$

where $T_b(\mathbf{x})$ is the proportion of deaths in the leaf of tree b that contains $\mathbf{x}$. The tuned RF used $B = 500$ and a grid over the number of predictors per split, the minimum leaf size and the maximum depth.

XGBoost. XGBoost is an additive tree model on the log-odds scale, $\operatorname{logit}\hat{p}(\mathbf{x}) = \sum_{k=1}^{K} f_k(\mathbf{x})$, fitted by gradient boosting with shrinkage (learning rate η) to minimise $\sum_i \ell_i + \sum_k \Omega(f_k)$ with

$$\Omega(f) = \gamma T + \tfrac{1}{2}\lambda\|\mathbf{w}\|^2,$$

where T is the number of leaves and $\mathbf{w}$ the leaf weights (Chen & Guestrin, 2016). γ and λ were kept at $\gamma = 0$ and the library default $\lambda = 1$; the number of trees, depth, learning rate, column and row subsampling and minimum child weight were tuned.

Support vector machine. The SVM solves the soft-margin problem $\min_{\mathbf{w},b,\boldsymbol{\xi}} \frac{1}{2}\|\mathbf{w}\|^2 + C\sum_i \xi_i$ subject to $\tilde{y}_i(\mathbf{w}^\top \phi(\mathbf{x}_i) + b) \geq 1 - \xi_i$ and $\xi_i \geq 0$, with $\tilde{y}_i \in \{-1, +1\}$ and the radial basis function kernel $K(\mathbf{x}, \mathbf{x}') = \exp(-\gamma\|\mathbf{x} - \mathbf{x}'\|^2)$ (Cortes & Vapnik, 1995). The SVM outputs a decision value $f(\mathbf{x})$; probabilities were obtained by Platt scaling, $\hat{p} = 1/(1 + \exp(Af(\mathbf{x}) + B))$, with A and B fitted by internal cross-validation (`probability=True`) (Platt, 1999). During tuning, AUROC was computed from the decision values, and the selected configuration was then refitted with Platt scaling. C and γ were tuned.

Two configurations per family. Each family was run (a) with the hyperparameters reported by the benchmark and (b) tuned by grid search inside the training data (stratified 5-fold cross-validation, AUROC as the criterion). Together with the unpenalised LR this gives nine configurations. RF was run in two benchmark versions, one with the literal parameters and one with R `randomForest` defaults (nine fully grown trees), because the benchmark lists R and Python parameter names together and the two readings imply very different trees. All grids are listed in Supplementary S9.

2.5 Evaluation framework (I5)

Repeated nested cross-validation (main estimate). Each training set was split into 10 stratified outer folds, and the split was repeated 5 times with different seeds, giving $J = 50$ outer folds. Inside every outer training fold, the 20% missingness rule, imputation, standardisation, LASSO (inner stratified 10-fold cross-validation, λ.1se) and hyperparameter tuning (inner stratified 5-fold cross-validation) were run in turn; the model was then refitted on the whole outer training fold and evaluated on the untouched outer validation fold (Figure 2). This estimates the performance of the whole modelling procedure, including variable selection and tuning, rather than of one fitted model (Cawley & Talbot, 2010; Varma & Simon, 2006). Preprocessing and LASSO were fitted once per outer training fold and were not refitted inside the inner tuning folds; outer validation observations were not used in preprocessing, variable selection or tuning. AUROC is reported as the mean $\pm$ SD over the 50 outer folds. In each repeat the 10 outer validation folds were pooled into out-of-fold (OOF) predictions covering the whole training set; AUPRC, Brier score, calibration, decision curves and thresholds were computed from these and averaged over the 5 repeats.

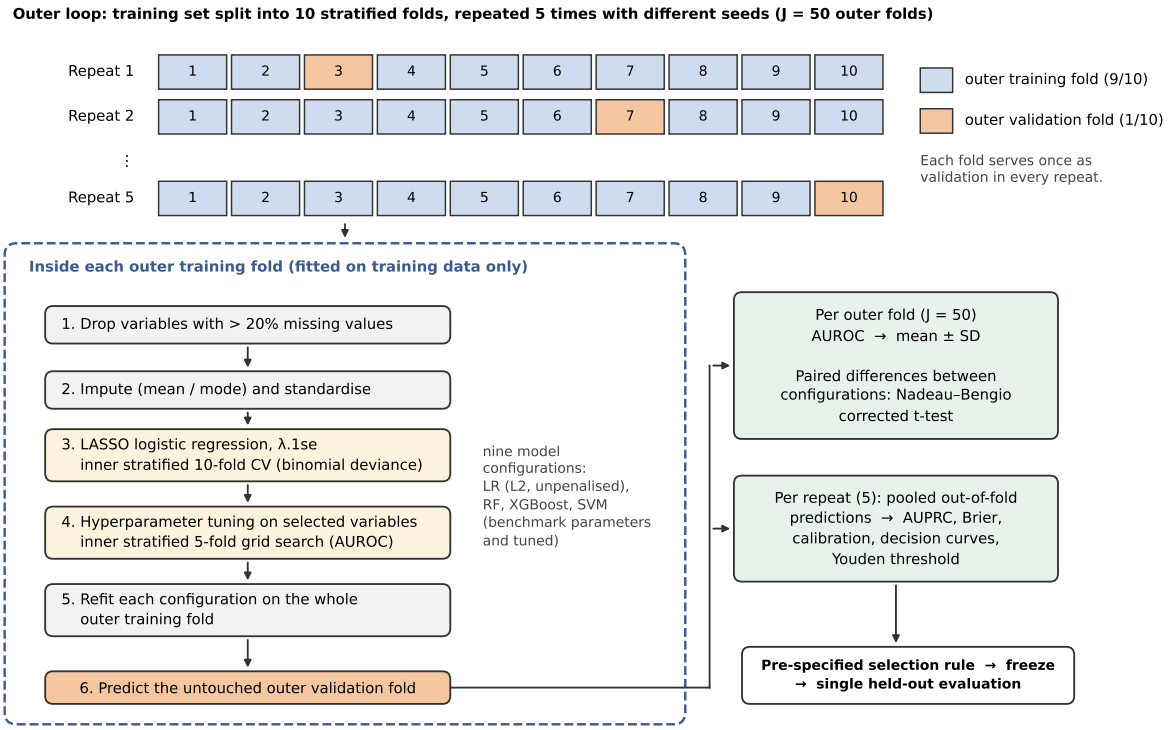


Figure 2: Repeated nested cross-validation. Every data-dependent step, including variable selection and tuning, is fitted inside each outer training fold and evaluated on the untouched outer validation fold.

Model comparison. Two models evaluated on the same folds were compared with the Nadeau–Bengio corrected resampled t-test (Nadeau & Bengio, 2003). With d_j the AUROC difference in outer fold j, $\bar{d}$ its mean and $\hat{\sigma}_d^2$ the sample variance of the J differences,

$$t = \frac{\bar{d}}{\sqrt{\left(\frac{1}{J} + \frac{n_{\text{test}}}{n_{\text{train}}}\right)\hat{\sigma}_d^2}},$$

referred to a t distribution with $J-1=49$ degrees of freedom, where $n_{\text{test}}/n_{\text{train}}$ is the ratio of validation-fold to training-fold size. The term $n_{\text{test}}/n_{\text{train}}$ inflates the variance to account for the overlap between training sets, which the naive paired t-test (variance $\hat{\sigma}_d^2/J$) ignores. The 95% confidence interval for the difference is $\bar{d} \pm t_{0.975,J-1}$ times the corrected standard error. In a simulation with 200 null datasets and the same 10-fold, 5-repeat design, the corrected test, which is an approximation, rejected the null hypothesis in 4.0% of datasets and the uncorrected paired t-test in 38.0%.

Pre-specified model-selection rule. (1) The configuration with the highest mean outer AUROC is the reference M*. (2) A configuration is retained as a candidate if its AUROC is lower than that of M* by at most 0.01, or by at most 0.02 with a 95% confidence interval for the difference that includes zero. (3) If the candidates include an LR, an LR is chosen, for parsimony and interpretability. (4) If both LRs are candidates, the one whose pooled-OOF calibration slope is closer to 1 is chosen, measured by the mean over repeats of |ln slope| (the unpenalised LR on a tie). (5) The held-out set plays no part. The rule was written before any landmark analysis was run; for the reference cohort it was written after exploratory

analyses and is therefore not blind.

Freezing and the held-out evaluation. After selection, the final landmark model was fitted on the whole training set with the same pipeline (LASSO with inner 10-fold cross-validation, then C chosen by inner 5-fold cross-validation). Its specification, the classification threshold and identifiers of the training data and code were recorded under version control before the held-out set was read. The held-out evaluation was run exactly once and was not repeated or followed by changes to the model. A tuned RF, frozen in the same way, served as a comparator. In the reference cohort the held-out set had been used in earlier exploratory model comparisons, so its results are reported only as a secondary estimate.

Held-out uncertainty. On held-out sets, the AUROC 95% confidence interval used the DeLong variance (DeLong et al., 1988), computed with the fast algorithm of Sun and Xu (2014), and the two frozen models were compared with the paired DeLong test. AUPRC confidence intervals came from a stratified bootstrap with 2,000 resamples (percentile method). Calibration intercepts and slopes were given Wald 95% confidence intervals from the recalibration models; these were computed after the confirmatory evaluation, from the same frozen predictions.

2.6 Metrics (I7)

Death was the positive class throughout.

Discrimination. AUROC, and AUPRC computed as average precision. The AUPRC of an uninformative model equals the prevalence $\bar{y}$, which is reported as its baseline.

Overall accuracy. The Brier score and the scaled Brier score (Steyerberg et al., 2010),

$$\text{Brier} = \frac{1}{n}\sum_{i=1}^{n}(\hat{p}_i - y_i)^2, \qquad \text{scaled Brier} = 1 - \frac{\text{Brier}}{\bar{y}(1-\bar{y})},$$

where $\bar{y}(1-\bar{y})$ is the Brier score of a model that predicts the prevalence for everyone, computed on the same data.

Calibration (Van Calster et al., 2019). A logistic recalibration model

$$\operatorname{logit} P(y=1) = a + b \operatorname{logit} \hat{p}$$

was fitted to the predictions. The calibration slope is b from fitting a and b together; the calibration intercept (calibration-in-the-large) is a from refitting with b fixed at 1, that is, with $\operatorname{logit}\hat{p}$ as an offset. A well-calibrated model has $a = 0$ and $b = 1$; $b < 1$ means predictions that are too extreme and $b > 1$ predictions that are too close to the mean.

Clinical utility. Decision-curve analysis (Vickers & Elkin, 2006) compares the net benefit of acting on the model at threshold probability p_t,

$$\text{NB}(p_t) = \frac{TP}{n} - \frac{FP}{n} \cdot \frac{p_t}{1-p_t},$$

with treating everyone, $\mathrm{NB}_{\mathrm{all}}(p_t) = \bar{y} - (1-\bar{y})\,\frac{p_t}{1-p_t}$, and treating no one (NB = 0). We evaluated the 56 thresholds from 0.05 to 0.60 in steps of 0.01 and counted the thresholds at which the model's net benefit exceeded both alternatives.

Classification. The threshold maximised Youden's index (Youden, 1950), $J = \mathrm{Se} + \mathrm{Sp} - 1$, on the pooled OOF predictions of each repeat; the median over the 5 repeats was carried unchanged to the held-out set. We report sensitivity, specificity, precision, F1 and the share of patients flagged as high risk.

2.7 Interpretation and sensitivity analyses

SHAP. For an LR on standardised predictors, the SHAP value of predictor j (under feature independence) is

$$\phi_j(\mathbf{x}) = \beta_j\,(x_j - \mathrm{E}[x_j]),$$

on the log-odds scale, with the expectation taken over the training data used as background (Lundberg & Lee, 2017). Because standardisation was fitted on the same training data, $\mathrm{E}[x_j] = 0$ and $\phi_j(\mathbf{x}) = \beta_j x_j$. Global importance is the mean $|\phi_j|$ over the held-out set. Tree models were explained with TreeExplainer.

Odds ratios. To describe associations, the selected predictors were refitted by unpenalised LR on the training set. These odds ratios differ from the penalised coefficients of the final model, and their Wald intervals do not account for variable selection, so they are exploratory. Because predictors are standardised, $\exp(\beta_j)$ is the odds ratio per standard deviation, with Wald 95% confidence interval $\exp(\beta_j \pm 1.96\,\mathrm{SE}_j)$; for binary predictors the coefficient was rescaled by the predictor's standard deviation to give the odds ratio for yes versus no. Collinearity was checked with variance inflation factors (VIF).

Sensitivity analyses. All were run with the full nested cross-validation on the same outer splits as the main analysis, and compared with it by the Nadeau–Bengio test: (i) reference cohort, comorbidities and CCI recomputed from prior plus chronic index-admission codes (I3); (ii) reference cohort, exclusion of patients admitted in 2020–2022, who are not covered by the benchmark's data release (new single 80/20 split); (iii) landmark cohort, comorbidities from prior admissions only (*prior-only*, the most conservative deployment scenario; I3); and (iv) landmark cohort, removal of all white-cell differential variables (*no-WBC-differential*; I2).

2.8 Exploratory extensions

The project brief proposed temporal features, comorbidity representations and a stacking ensemble. These were evaluated after the main analysis had been completed, on the landmark development set only (2,045 patients, 544 deaths); the held-out set was not used. Three additions were made to the 56 static candidates. (i) *First-day trajectories*: five statistics (first-to-last change, hourly slope, standard deviation, late-minus-early mean and number of measurement times) for each of 25 laboratory and vital-sign series recorded 0–24 h after ICU admission, giving 125 candidates. (ii) *Prior-diagnosis burden*: four variables describing earlier admissions and their ICD codes. (iii) *Diagnosis embedding*: each patient's ICD codes from admissions that ended before the index admission were converted to a TF-IDF document and compressed to 32 dimensions by truncated singular value decomposition, both fitted within training folds. Stacking

combined LR, RF and XGBoost base learners through an L2 LR meta-model trained on their internal 3-fold out-of-fold predictions.

To keep the computation feasible, all extension models and their static controls used one fixed-parameter recipe (LR with $C = 1$; RF with 200 trees of maximum depth 10; XGBoost with 100 trees of depth 2; LASSO with 5 internal folds and 30 values of λ) on the same 5×10 outer splits as the main analysis, with missingness screening, imputation, LASSO and the embedding refitted within each training fold. Each extension was compared with its matched control on the same folds by the corrected test, with Holm adjustment over 15 comparisons per metric. Because this recipe omits inner tuning, the extensions are compared only with their own static controls, not with Table 4. Input features were verified against our own extraction (Supplementary S12).

3 Results

3.1 Cohort: who the benchmark's criteria exclude (I1, I2)

Of 8,992 patients with a DN code, 3,189 had a DN-coded admission with an ICU stay (Figure 1). The landmark cohort comprised 2,557 patients with 680 deaths within one year (26.6%): 594 patients whose first ICU stay ended within 24 h and 38 who died within 24 h were excluded. The reference cohort, built with the benchmark's criteria, comprised 1,960 patients with 494 deaths (25.2%).

The retention at each benchmark step closely matched the benchmark (hospital stay > 48 h, 95.5% vs 95.4%; complete blood count, 64.3% vs 64.2%), and the larger cohort (1,960 vs 1,313 in the benchmark) probably reflects the extended data release (12.6% of the reference cohort was admitted in 2020–2022). Exclusion at the blood-count step was almost entirely due to a missing white-cell differential (35.6% of candidates); the core blood count was missing in only 1%. The landmark cohort keeps these patients: 873 (34.1%) had no differential, and their one-year mortality was 28.6% (250/873), against 25.5% (430/1,684) for patients with a differential (χ^2 p = 0.09). In total, 1,452 landmark patients (56.8%) had a prior admission from which comorbidities could be taken.

Non-survivors were older and had higher RDW, CCI and severity scores and lower chloride (Table 2).

Table 2: Selected baseline characteristics of the landmark cohort by one-year outcome (median [IQR] or n [%]; observed values before imputation). Full tables for both cohorts are in Supplementary S3.

Variable	Overall (n = 2,557)	Survivors (n = 1,877)	Non-survivors (n = 680)	p
Age, years	67 (58–75)	65 (57–73)	71 (63–79)	<0.001
Male sex	1,520 (59.44%)	1,097 (58.44%)	423 (62.21%)	0.087
Congestive heart failure	1,178 (46.07%)	775 (41.29%)	403 (59.26%)	<0.001
Renal disease	1,299 (50.80%)	877 (46.72%)	422 (62.06%)	<0.001
Obesity (coded)	859 (33.59%)	664 (35.38%)	195 (28.68%)	0.002
CCI (without age)	5 (4–7)	5 (3–6)	6 (5–8)	<0.001
RDW, %	14.65 (13.60–16.20)	14.30 (13.40–15.65)	15.81 (14.49–17.63)	<0.001
Chloride (max), mmol/L	105 (100–109)	106 (101–109)	103 (98–108)	<0.001
Anion gap (min), mmol/L	13 (10–15)	12 (10–14)	14 (11–17)	<0.001
Urea nitrogen (min), mg/dL	23 (15–39)	21 (14–34)	32 (20–53)	<0.001
Lymphocytes, 10^9/L †	1.27 (0.79–1.95)	1.39 (0.87–2.08)	0.95 (0.60–1.47)	<0.001

Variable	Overall (n = 2,557)	Survivors (n = 1,877)	Non-survivors (n = 680)	p
Respiratory rate, /min	18.28 (16.43–20.62)	18.08 (16.32–20.20)	19.12 (16.99–21.75)	<0.001
SAPS-II	36 (29–45)	34 (27–42)	42 (35–52.25)	<0.001
APS-III	45 (35–57)	42 (34–53)	54 (43–67)	<0.001
OASIS	30 (25–36)	30 (24–35)	33.5 (27–40)	<0.001
No white-cell differential	873 (34.14%)	623 (33.19%)	250 (36.76%)	0.092
Prior admission	1,452 (56.79%)	1,028 (54.77%)	424 (62.35%)	<0.001

† 34.1% missing (no differential performed).

3.2 Definition errors in the benchmark (I4)

The comparisons in this section use the reference cohort.

Mortality. One-year mortality from ICU admission was 25.2% in the reference cohort, against 38.2% in the benchmark. In-hospital mortality was only 7.1%, so the benchmark's figure cannot be in-hospital mortality, although its abstract names in-hospital death. Counting every recorded date of death regardless of timing gave 42.8%, the same order as the benchmark, and 49% of patients with an out-of-hospital date of death died more than 365 days after ICU admission. The benchmark may therefore have used an outcome without a fixed follow-up window; because it does not state its definition, this cannot be confirmed.

White-cell differential units. The benchmark reported lymphocytes of 3.69 (1.27–118.32), neutrophils of 23.90 (8.63–831.30) and monocytes of 2.00 (0.60–43.64) $\times$ 10^9/L, which are physiologically implausible. With absolute counts we obtained 1.24 (0.77–1.91), 8.59 (5.99–12.57) and 0.55 (0.35–0.84). The pattern of the derived ratios indicates mixing of scales across patients. Ratios that are unaffected when a patient's three cell counts share one scale agreed with the benchmark (NLR 6.82 vs 6.80; LMR 2.31 vs 2.42; SII 1,336 vs 1,342), whereas ratios that combine platelets with a differential count did not (PLR 159 vs 36; PNR 22 vs 6.6). When we deliberately mixed all differential items, PLR and PNR fell to 24.00 and 3.67, moving towards the benchmark's values. With correct units, neutrophils were selected by LASSO only at λ.min, not at λ.1se, and were not among the five most important predictors, whereas they ranked second in the benchmark.

Hypertension. Hypertension prevalence was 84.6% under the full Quan–Elixhauser definition, against 37.3% in the benchmark; counting only uncomplicated hypertension codes gave 35.4%, close to the benchmark. The narrow definition reverses the apparent effect. Under the full definition, hypertension was unrelated to one-year death (unadjusted OR 1.07, 95% CI 0.80–1.42). Under the uncomplicated definition it appeared protective (OR 0.49, 0.39–0.62), because the 964 patients who had only complicated codes, and were therefore coded as normotensive, had a one-year mortality of 31.3% (OR 1.91, 1.55–2.35). LASSO did not select hypertension under either definition, and the selected variables were identical, so the models were unaffected.

3.3 The benchmark pipeline under nested cross-validation (I5, I6)

LASSO on the reference training set selected 12 variables at λ.1se (λ = 0.02667, against 0.02687 in the benchmark), seven of which overlapped with the benchmark's 12 (age, anion gap, RDW, chloride, PT, APS-III, CCI); at λ.min the overlap was 11 of 12.

Under nested cross-validation, six configurations met the selection criteria, with AUROC 0.786–0.791 (Table 3; full results in Supplementary S2). The benchmark's two RF configurations and its SVM were significantly worse, with confidence intervals for the difference that excluded zero. Among the benchmark-parameter models, only XGBoost matched the best model.

Table 3: Reference cohort: nested cross-validation on the training set (1,568 patients; 5 × 10 folds) and the benchmark's reported validation AUCs.

Model	Benchmark validation AUC (Huang et al., 2025)	Nested-CV AUROC (mean ± SD)	AUPRC	Brier	Calibration slope	ΔAUROC vs best (95% CI)
XGBoost (benchmark parameters)	0.732	**0.791 ± 0.035**	0.559	0.150	0.96	reference
XGBoost (tuned)	—	0.790 ± 0.035	0.558	0.150	0.96	0.000 (−0.005, 0.006)
LR (L2, tuned)	—	0.788 ± 0.036	0.536	0.154	1.19	0.002 (−0.009, 0.014)
RF (tuned)	—	0.788 ± 0.036	0.548	0.154	1.48	0.002 (−0.006, 0.011)
SVM (tuned)	—	0.788 ± 0.036	0.541	0.153	0.90	0.003 (−0.008, 0.014)
LR (unpenalised)	0.773	0.786 ± 0.035	0.550	0.151	0.85	0.005 (−0.007, 0.016)
RF (benchmark, literal parameters)	0.780	0.767 ± 0.038	0.519	0.156	0.94	0.024 (0.006, 0.041)
RF (benchmark, R defaults)	0.780	0.724 ± 0.039	0.440	0.171	0.15	0.066 (0.043, 0.090)
SVM (benchmark parameters)	0.679	0.723 ± 0.044	0.497	0.163	0.96	0.068 (0.042, 0.093)

AUPRC baseline = 0.252. The first six rows met the selection criteria of the pre-specified rule, which selected the unpenalised LR (for this cohort the rule was not blind; Section 2.5). On the secondary held-out set (392 patients), the unpenalised LR had AUROC 0.803 (0.753–0.853) and calibration slope 0.96.

Three further results bear on the benchmark's pipeline.

- **RF overfitting (I6).** With R defaults (nine fully grown trees), the RF reproduced the benchmark's training AUC of 0.999 but had a validation AUC of 0.756 on a single split, a nested-CV AUROC of 0.724 and a calibration slope of 0.15, indicating severely overconfident predictions. With the literal parameters the training AUC was only 0.887. The benchmark's 0.999 is therefore consistent with fully grown trees, although other causes cannot be ruled out.
- **Variable selection outside cross-validation (I5).** Treating the LASSO selection as fixed, as the

benchmark did, overstated the cross-validated AUROC of the unpenalised LR by 0.011 (0.793 vs 0.781), and of the nine configurations by 0.008–0.030 (single repeat of 10-fold cross-validation).

- **Preprocessing leakage (I5).** Fitting imputation and standardisation on all data, as in the benchmark, did not change the LASSO selection and changed model metrics by at most 0.001. Missingness was at most 7.04%, which probably explains why this form of leakage had negligible impact here.

A single validation set could not rank the models: on the 392-patient validation set, no pair of tuned models differed significantly after Holm correction (smallest adjusted p = 0.13).

3.4 Improved landmark model (main result)

Nested cross-validation. LASSO selected 13–23 of the 56 candidates per outer fold (median 17). Age, anion gap, CCI, OASIS, APS-III, respiratory rate, RDW, maximum chloride and obesity were selected in all 50 folds; `wbc_diff_missing` was selected in 13 of 50. Of the six comorbidities, only obesity was ever selected. Six configurations again met the selection criteria (AUROC 0.794–0.801; Table 4), and the benchmark's RF and SVM configurations were again not.

Table 4: Landmark cohort: nested cross-validation on the training set (2,045 patients; 5 × 10 folds).

Model	AUROC (mean ± SD)	AUPRC	Brier	Calibration slope	ΔAUROC vs best (95% CI)
LR (L2, tuned)	**0.801 ± 0.026**	0.594	0.151	1.11	reference
LR (unpenalised)	0.799 ± 0.026	0.592	0.151	0.90	0.003 (−0.001, 0.006)
SVM (tuned)	0.798 ± 0.028	0.584	0.153	0.92	0.004 (−0.002, 0.009)
RF (tuned)	0.797 ± 0.028	0.591	0.155	1.46	0.004 (−0.005, 0.014)
XGBoost (tuned)	0.795 ± 0.028	0.596	0.152	0.99	0.006 (−0.004, 0.015)
XGBoost (benchmark parameters)	0.794 ± 0.029	0.597	0.152	1.02	0.008 (−0.002, 0.017)
RF (benchmark, literal parameters)	0.781 ± 0.029	0.560	0.157	1.09	0.020 (0.008, 0.033)
SVM (benchmark parameters)	0.756 ± 0.035	0.550	0.161	0.98	0.046 (0.027, 0.065)
RF (benchmark, R defaults)	0.740 ± 0.037	0.488	0.170	0.18	0.061 (0.039, 0.083)

AUPRC baseline = 0.266.

Model selection. The pre-specified rule selected the L2-penalised LR on calibration (mean |ln slope| 0.102 vs 0.109), although its discrimination was similar to that of the unpenalised LR (difference 0.003). Inner tuning chose C = 0.01 in all 50 folds.

Final model. Refitted on the whole training set (544 events), LASSO selected 22 variables (λ.1se = 0.0168) and inner tuning chose C = 0.01. The frozen model and the tuned RF comparator were evaluated once on the held-out set (512 patients, 136 deaths; Table 5, Figure 3).

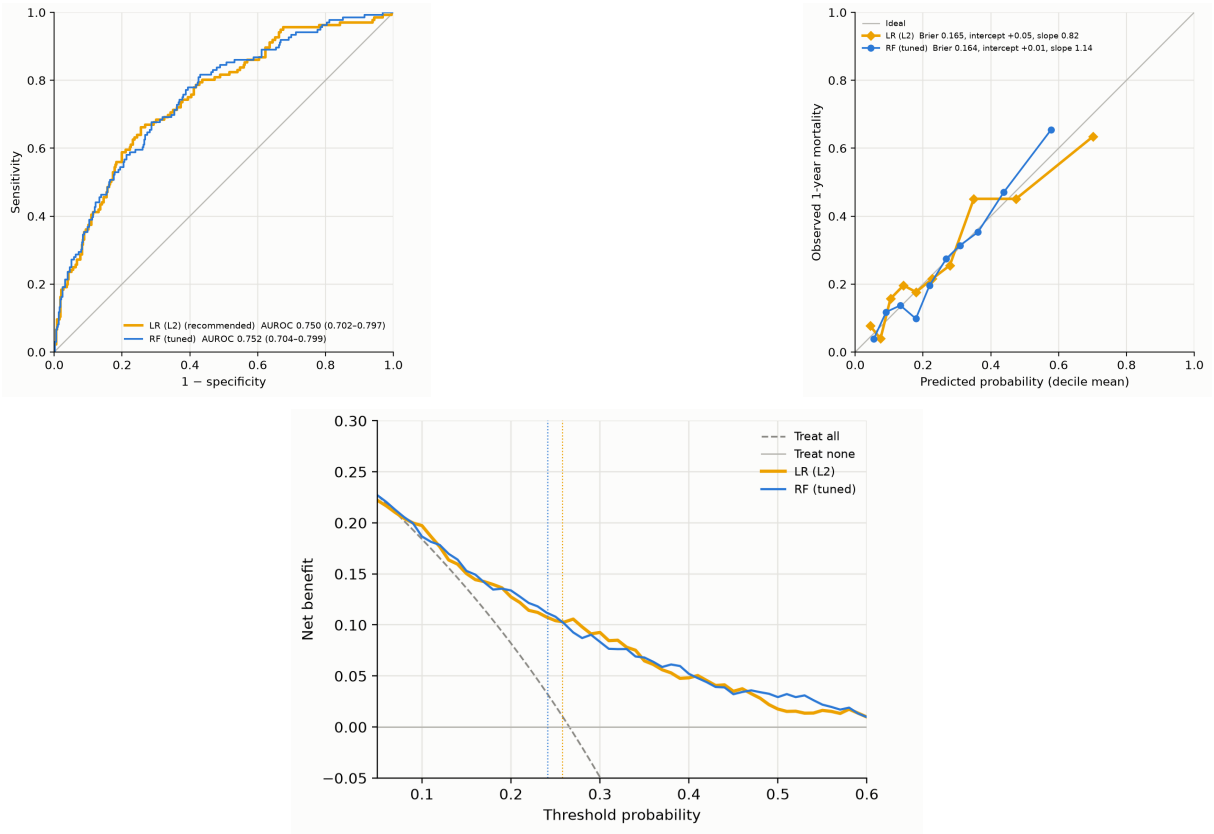


Figure 3: Frozen landmark models on the held-out set (512 patients, 136 deaths): (a) ROC curves, (b) calibration, (c) decision curves, with dotted vertical lines marking each model's frozen Youden threshold.

Table 5: Landmark cohort: confirmatory held-out evaluation of the frozen models.

Metric	LR (L2), selected	RF (tuned), comparator
AUROC (95% CI)	0.750 (0.702–0.797)	0.752 (0.704–0.799)
AUPRC (95% CI; baseline 0.266)	0.533 (0.459–0.611)	0.529 (0.460–0.615)
Brier / scaled Brier	0.165 / 0.153	0.164 / 0.161
Calibration intercept (95% CI)	+0.049 (−0.171, 0.270)	+0.013 (−0.197, 0.223)
Calibration slope (95% CI)	0.82 (0.62–1.03)	1.14 (0.86–1.42)
Threshold (OOF Youden)	0.258	0.241
Sensitivity / specificity	0.684 / 0.697	0.779 / 0.596
Precision / F1	0.449 / 0.542	0.411 / 0.538
Flagged as high risk	40.4%	50.4%

Paired DeLong test, LR − RF: −0.002 (−0.021, 0.016), p = 0.83.

Held-out estimate below the cross-validated estimate. The held-out AUROC was lower than the nested-CV estimate for both LR (0.750 vs 0.801) and RF (0.752 vs 0.797). The available analyses did not establish the cause of this difference; post-hoc diagnostic analyses are reported in Supplementary S13.

3.5 Calibration and clinical utility (I7)

Calibration. Out of fold, the selected L2 LR had a calibration slope of 1.11 (predictions slightly too close to the mean) and the unpenalised LR 0.90. On the held-out set, the L2 LR's slope was 0.82 (95%

CI 0.62–1.03) with intercept +0.049 (−0.171 to 0.270): the point estimate suggests slightly too extreme predictions, the opposite direction to the cross-validated estimate, but with 136 deaths neither value differed clearly from perfect calibration (p = 0.09 for slope = 1). The tuned RF had out-of-fold slopes of 1.46 in the landmark cohort and 1.48 in the reference cohort (predictions shrunk towards the mean), and 1.14 (0.86–1.42) on the held-out set. The benchmark's R-default RF had slopes of 0.18 and 0.15. Discrimination alone would not have revealed these differences, and the change in the selected model's slope between samples means that calibration should be assessed in any target population.

Overall accuracy. The held-out Brier score of the L2 LR was 0.165 (scaled Brier 0.153), similar to the RF comparator (0.164; 0.161).

Decision curves. On the held-out set, the L2 LR had net benefit above both treat-all and treat-none at 53 of 56 thresholds (0.08–0.60), and the tuned RF at 55 of 56 (0.06–0.60). In the reference cohort, the unpenalised LR was above both at 56 of 56 thresholds in nested cross-validation and on its held-out set.

Classification. At its out-of-fold Youden threshold (0.258), the landmark LR flagged 40.4% of held-out patients as high risk, of whom 44.9% died within one year (precision 0.449), with sensitivity 0.684 and specificity 0.697. In the reference cohort, a threshold of 0.5 flagged 16.3% of patients and yielded a sensitivity of 0.44. The benchmark reported precision 0.735 and recall 0.858 for its RF; these values are internally consistent only if survival was the positive class, so they are not comparable with ours.

3.6 What each improvement changed (sensitivity analyses)

Table 6: Sensitivity analyses (nested cross-validation on the same outer splits; ΔAUROC = sensitivity − main, Nadeau–Bengio 95% CI).

Cohort	Analysis	Model	AUROC	ΔAUROC (95% CI)
Reference	Comorbidity timing	LR (unpenalised)	0.787	+0.001 (−0.006, 0.008)
Reference	Comorbidity timing	RF (tuned)	0.790	+0.002 (−0.006, 0.010)
Landmark	Prior-admission comorbidities only	LR (L2)	0.787	−0.014 (−0.023, −0.005)
Landmark	Prior-admission comorbidities only	RF (tuned)	0.789	−0.008 (−0.016, −0.0001)
Landmark	No white-cell differential	LR (L2)	0.802	+0.001 (−0.005, 0.007)
Landmark	No white-cell differential	RF (tuned)	0.795	−0.002 (−0.010, 0.005)

Comorbidity timing (I3). In the reference cohort, recomputing comorbidities from prior plus chronic index-admission codes changed the CCI for 843 patients, and 304 of 611 myocardial-infarction diagnoses (about half) came only from acute index-admission codes. Discrimination was nevertheless unchanged, and CCI was still selected in every fold. In the landmark cohort, using prior admissions only lowered AUROC by 0.014 for the LR, with a confidence interval excluding zero. The largest change was in obesity, whose selection fell from 50 to 4 of 50 folds. This prior-only estimate (about 0.787) is the most conservative deployment scenario, because it does not assume that chronic index-admission codes were present at admission.

White-cell differential (I2). Removing the eight differential variables and the missingness indicator did not change discrimination, AUPRC, Brier score or calibration, and reduced the median number of selected

variables from 17 to 13. The benefit of I2 is therefore in the cohort it allows, not in discrimination.

Data release. Excluding the 247 patients admitted in 2020–2022 from the reference cohort (new single 80/20 split) gave validation AUROCs of 0.799 for the unpenalised LR and 0.802 for the tuned RF, against 0.803 and 0.793 in the main split, and RDW and CCI remained the two most important predictors of the LR.

3.7 Predictors and interpretation

RDW and CCI ranked first and second in every model, and age ranked third in both LRs; the tuned RF ranked SAPS-II third (Table 7; Figure 4). In the landmark model (unpenalised refit, odds ratio per SD), higher risk was associated with RDW (OR 1.67, 95% CI 1.47–1.89), age (1.50, 1.30–1.75), CCI (1.44, 1.27–1.63), APS-III (1.38, 1.14–1.67), OASIS (1.23, 1.04–1.46), respiratory rate (1.20, 1.06–1.34) and minimum calcium (1.16, 1.02–1.32). Lower risk was associated with higher maximum chloride (0.82, 0.73–0.94) and coded obesity (yes vs no, 0.63, 0.49–0.81). Male sex (1.31, 1.03–1.66) and a missing white-cell differential (1.33, 1.04–1.71) were associated with higher risk. All variance inflation factors were below 5 (largest 4.74, SAPS-II), and SAPS-II, LODS, NLR and SII were not significant once the other scores and ratios were included.

Table 7: Top-5 predictors by mean |SHAP|.

Rank	Benchmark (RF) (Huang et al., 2025)	Reference: LR (unpenalised)	Reference: RF (tuned)	Landmark: LR (L2)
1	RDW	RDW	RDW	RDW
2	Neutrophils	CCI	CCI	CCI
3	CCI	Age	SAPS-II	Age
4	Chloride (max)	APS-III	APS-III	APS-III
5	PT (min)	OASIS	Chloride (max)	Obesity

3.8 Exploratory extensions

All extension results come from the development set.

None of the extensions clearly improved on the static LR (Table 8). Stacking had the highest AUROC, 0.805 against 0.799, a difference of 0.005 (95% CI −0.001 to 0.011; p = 0.077, Holm-adjusted p = 1.00). Its direction was consistent, with higher AUROC in 41 of 50 folds and a higher pooled out-of-fold AUROC in all five repeats, and its out-of-fold calibration slope (1.05) was closer to 1 than that of LR (0.90); its net benefit, however, was higher at only 31 of 56 thresholds. Trajectory features entered every fit (median 4.5 selected per fold, range 1–13; SpO2 variability in 48 of 50 folds) but lowered AUROC slightly for LR, RF and XGBoost (by 0.004, 0.001 and 0.002). Prior diagnoses were available for 1,183 development patients (57.8%). Neither the burden variables nor the embedding improved LR, and the embedding's small gain for RF (0.002 over the static RF) came with a worse Brier score and a calibration slope further from 1 (1.35). Adding all extensions to the stacking ensemble changed AUROC by less than 0.001. All comparisons had Holm-adjusted p = 1.00; the full results for 14 models are in Supplementary S12.

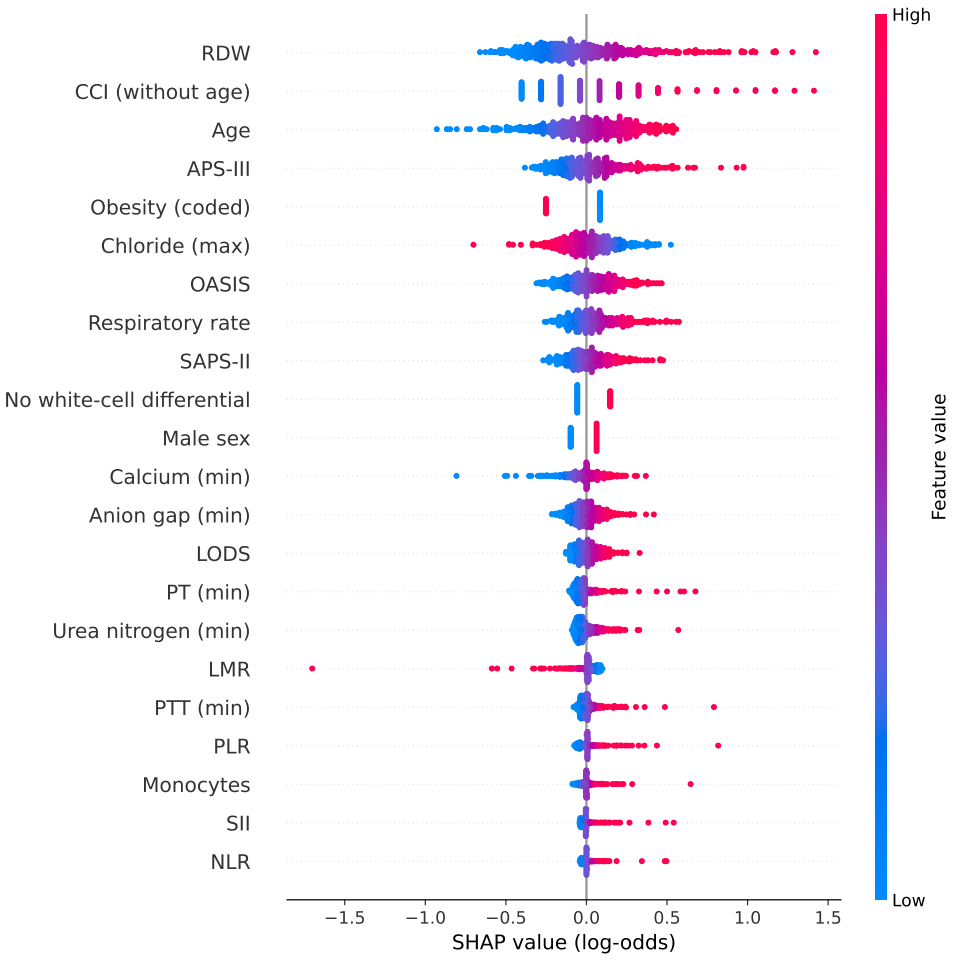


Figure 4: SHAP summary of the frozen landmark LR (L2) on the held-out set (512 patients), log-odds scale, predictors ordered by mean |SHAP|. Colour shows the predictor value. For patients without a white-cell differential, the differential-derived variables take their imputed training means and appear as a cluster near zero. The corresponding plots for the reference cohort are in Supplementary S2.

Table 8: Exploratory extensions on the landmark development set (fixed-parameter recipe, 5 × 10 folds; ΔAUROC against the matched reference on the same folds, with Nadeau–Bengio-corrected 95% CI).

Model	AUROC (mean ± SD)	AUPRC	Brier	Reference	ΔAUROC (95% CI)
Static LR (L2)	0.799 ± 0.026	0.607	0.151	—	—
Static stacking	0.805 ± 0.027	0.616	0.150	Static LR	0.005 (−0.001, 0.011)
LR + trajectories	0.796 ± 0.026	0.602	0.153	Static LR	−0.004 (−0.010, 0.003)
LR + diagnosis burden	0.799 ± 0.027	0.607	0.151	Static LR	0.000 (−0.002, 0.002)
LR + burden + embedding	0.796 ± 0.027	0.606	0.152	LR + burden	−0.003 (−0.011, 0.004)
Stacking, all extensions	0.805 ± 0.027	0.616	0.150	Static stacking	0.000 (−0.008, 0.008)

4 Discussion

4.1 Principal findings

Starting from the benchmark's framework, we corrected seven issues and re-evaluated the result. The improved landmark model predicts one-year mortality at a defined time, 24 h after ICU admission, in patients alive and in the ICU at that time, and keeps the third of patients that the benchmark's blood-count criterion excluded. Its nested-CV AUROC was 0.801 on the development set and its AUROC on the held-out set was 0.750. Within the benchmark's framework of shared LASSO selection and our tuning grids, we found no evidence that RF, XGBoost or SVM outperforms LR; because the trees received only LASSO-selected variables, this does not show that LR matches tree models in general. The benchmark's conclusion that RF is superior was not reproduced: its RF configuration was significantly worse than the best model, and its training AUC of 0.999 is reproduced by a forest of nine fully grown trees that memorises the training data. The leading predictors were stable across cohorts and models: RDW, comorbidity burden, age, illness severity and low chloride.

4.2 Impact of each improvement

Table 9 summarises what each modification changed. Some changed the conclusions; others changed little and are reported for completeness.

Table 9: Observed effect of each improvement.

Improvement	Observed effect
I1 Landmark at 24 h	Gives the model a defined time of use and removes a criterion that depends on the future. The landmark model's nested-CV AUROC (0.801) was similar to that of the reference cohort (0.786–0.791), although the cohorts differ and this is not a controlled comparison.
I2 Keep patients without a differential	Retained 873 patients (34.1%), whose mortality was 28.6% against 25.5% (p = 0.09); the missingness indicator had OR 1.33 (1.04–1.71) but was selected in only 13 of 50 folds. Removing all differential variables changed AUROC by +0.001 (−0.005, 0.007): no effect on discrimination.
I3 Comorbidity timing	Changed the CCI for 843 reference patients; about half of the myocardial-infarction diagnoses (304 of 611) rested only on acute codes. AUROC did not change (+0.001). Relying on prior admissions only lowered AUROC by 0.014 and removed obesity from the model (50 to 4 of 50 folds), so the effect is mainly on interpretation.

Improvement	Observed effect
I4 Definitions	A fixed 365-day window gave 25.2% mortality, against 42.8% when any recorded death is counted. The full hypertension definition removed a spurious protective effect (OR 0.49 vs 1.07) without changing the LASSO selection. With absolute counts, neutrophils were not selected at λ.1se.
I5 Nested evaluation	Selecting variables outside cross-validation overstated AUROC by 0.008–0.030 in a single-repeat analysis (0.011 for the unpenalised LR). Preprocessing leakage changed metrics by at most 0.001: no material effect at this missingness. A single validation set could not separate the tuned models, and the held-out estimate was 0.05 below the cross-validated one.
I6 Tuning all families	The largest effect. The benchmark's RF configurations reached 0.724 and 0.767, significantly below the best model (0.791), whereas the tuned RF (0.788) met the selection criteria. With tuning, no advantage of RF was found.
I7 Calibration and decision curves	Revealed differences that discrimination hides: the R-default RF had slope 0.15, and the selected LR's slope moved from 1.11 out of fold to 0.82 on the held-out set, so calibration needs to be assessed in the target population. Net benefit exceeded both default strategies at 53 of 56 thresholds.

4.3 Comparison with the benchmark

We found no RF advantage: once all families were tuned inside the folds, the differences between them were within 0.005 in the reference cohort and 0.008 in the landmark cohort. The benchmark's reported advantage of RF rested on a single 263-patient validation set, on which a confidence interval spans about $\pm$0.05 (0.723–0.837 for its RF), and on a training AUC that reflects memorisation rather than performance, as the benchmark's authors themselves noted. Because the cohort, data release and outcome changed together, the difference cannot be attributed to any single correction.

The held-out AUROCs of the two studies (0.750 here, 0.780 in the benchmark) should not be read as a direct comparison, because their cohorts, outcome definitions and validation procedures differ. In our held-out set, LR and RF achieved AUROCs of 0.750 and 0.752.

4.4 Biological plausibility of the main predictors

RDW was the strongest predictor in every analysis, in agreement with the benchmark and with studies linking RDW to mortality in critically ill patients (Bazick et al., 2011). Elevated RDW reflects heterogeneous erythropoiesis driven by inflammation, oxidative stress, nutritional deficiency and renal dysfunction, which are prevalent in long-standing diabetes. **CCI** captures the cardiovascular and renal burden that typically accompanies DN (Feldman et al., 2019), and its importance persisted after recomputing it from chronic codes only. **Low chloride**, associated with higher risk in both our work and the benchmark, is a recognised marker of poor outcome in critical illness (Tani et al., 2012). The apparently protective effect of **coded obesity** was concentrated in patients whose obesity code appeared only during the index admission: their one-year mortality was 18.8% (n = 416), against 26.4% (n = 443) when obesity was coded in a prior admission and 28.6% (n = 1,698) without an obesity code. The association may partly reflect differences in documentation rather than a true "obesity paradox", and should be tested with measured body-mass index.

4.5 Limitations

- **Single-centre, retrospective data and no external validation.** The held-out set is an internal validation from the same centre and database, and has only 136 events, which limits the precision of the confirmatory estimate.
- **Information timing at the landmark.** DN itself was identified from ICD codes of the index admission, which are assigned at discharge, so real-time use requires DN status to be known from history or earlier admissions. The main estimate also assumes that chronic index-admission codes were known at 24 h (using prior admissions only lowers AUROC to about 0.787), and laboratory windows used measurement rather than reporting times. The removal of acute codes reduces, but does not eliminate, the use of information that may not be available at 24 h.
- **Calibration.** The slope was 1.11 out of fold and 0.82 on the held-out set, so calibration should be assessed, and the model recalibrated where needed, in any target population.
- **Only first-day features.** Treatments were not used, within-day dynamics were tested only in the exploratory extensions, and cause of death is unavailable.
- **Shared LASSO selection.** All model families received the same linearly selected variables, so predictors that act only through non-linear effects or interactions may have been removed before the tree models were fitted. Our model comparison applies to this framework, not to the model families in general.
- **Exploratory extensions.** They were run after the held-out results had been seen, with fixed rather than tuned hyperparameters, and were not evaluated on held-out data; the gain from stacking is uncertain (CI including zero) rather than absent.
- **Non-blind model selection in the reference cohort.** The selection rule was written after exploratory analyses of the reference cohort, and its held-out set had been used in those analyses; only the landmark analyses were fully pre-specified.

4.6 Future work

- **Temporal validation.** Train on earlier and test on later MIMIC-IV year groups, and validate externally in data with one-year follow-up.
- **Calibration and thresholds.** Assess calibration in the target population and choose thresholds for specific clinical actions, such as early goals-of-care discussion or targeted follow-up, according to their harms and costs.
- **Interpretation.** SHAP values describe the model, not causes of death, and divide credit among correlated predictors such as APS-III, OASIS, SAPS-II and CCI; interaction-aware or grouped explanations would help.
- **Information timing and model families.** Check laboratory reporting times, restrict inclusion to DN known before admission, and compare tree models fitted without LASSO preselection.
- **Objective measurements.** Use measured body-mass index to test the obesity finding.
- **Richer predictors and models.** Add HbA1c and treatment variables, and test stacking with tuned base learners on new data. Clinical notes, SHAP interaction values and counterfactual explanations, also proposed for this project, were not run.

5 Conclusion

We took a published machine-learning framework for one-year mortality in ICU patients with diabetic neuropathy and corrected its cohort definition, predictor definitions and evaluation. The resulting landmark model predicts at 24 h after ICU admission and was frozen before a single confirmatory evaluation. Within the benchmark's LASSO-based framework and under repeated nested cross-validation, we found no evidence that tree ensembles outperform LR; the benchmark's reported RF advantage is consistent with an overfitted configuration evaluated on a single small validation set. RDW, comorbidity burden and age were consistently the strongest predictors. Exploratory trajectory features and diagnosis embeddings did not improve discrimination, and the gain from stacking was uncertain. The landmark model had an AUROC of 0.801 in nested cross-validation and 0.750 on the held-out set; before clinical use, its calibration should be assessed locally, the availability of its inputs at 24 h confirmed, and its performance validated temporally or externally.

References

Bazick, H. S., Chang, D., Mahadevappa, K., Gibbons, F. K., & Christopher, K. B. (2011). Red cell distribution width and all-cause mortality in critically ill patients. *Critical Care Medicine*, *39*(8), 1913–1921. https://doi.org/10.1097/CCM.0b013e31821b85c6

Breiman, L. (2001). Random forests. *Machine Learning*, *45*, 5–32. https://doi.org/10.1023/A:1010933404324

Cawley, G. C., & Talbot, N. L. C. (2010). On over-fitting in model selection and subsequent selection bias in performance evaluation. *Journal of Machine Learning Research*, *11*, 2079–2107.

Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794. https://doi.org/10.1145/2939672.2939785

Collins, G. S., Reitsma, J. B., Altman, D. G., & Moons, K. G. M. (2015). Transparent reporting of a multivariable prediction model for individual prognosis or diagnosis (TRIPOD): The TRIPOD statement. *BMJ*, *350*, g7594. https://doi.org/10.1136/bmj.g7594

Cortes, C., & Vapnik, V. (1995). Support-vector networks. *Machine Learning*, *20*, 273–297. https://doi.org/10.1007/BF00994018

DeLong, E. R., DeLong, D. M., & Clarke-Pearson, D. L. (1988). Comparing the areas under two or more correlated receiver operating characteristic curves: A nonparametric approach. *Biometrics*, *44*(3), 837–845. https://doi.org/10.2307/2531595

Feldman, E. L., Callaghan, B. C., Pop-Busui, R., Zochodne, D. W., Wright, D. E., Bennett, D. L., Bril, V., Russell, J. W., & Viswanathan, V. (2019). Diabetic neuropathy. *Nature Reviews Disease Primers*, *5*, 41. https://doi.org/10.1038/s41572-019-0092-1

Friedman, J., Hastie, T., & Tibshirani, R. (2010). Regularization paths for generalized linear models via coordinate descent. *Journal of Statistical Software*, *33*(1), 1–22. https://doi.org/10.18637/jss.v033.i01

Huang, Y., Mo, W., Liu, Y., Wang, S., Zeng, Y., Jiang, P., Wang, H., & Bi, J. (2025). Machine learning models predict mortality risk in diabetic neuropathy patients using MIMIC-IV data. *Scientific Reports*, *15*, 38702. https://doi.org/10.1038/s41598-025-22363-x

Johnson, A. E. W., Bulgarelli, L., Shen, L., Gayles, A., Shammout, A., Horng, S., Pollard, T. J., Hao, S., Moody, B., Gow, B., Lehman, L.-w. H., Celi, L. A., & Mark, R. G. (2023). MIMIC-IV, a freely accessible electronic health record dataset. *Scientific Data*, *10*, 1. https://doi.org/10.1038/s41597-022-01899-x

Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, *30*, 4765–4774.

Nadeau, C., & Bengio, Y. (2003). Inference for the generalization error. *Machine Learning*, *52*, 239–281. https://doi.org/10.1023/A:1024068626366

Platt, J. C. (1999). Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In A. J. Smola, P. Bartlett, B. Schölkopf, & D. Schuurmans (Eds.), *Advances in large margin classifiers* (pp. 61–74). MIT Press.

Pop-Busui, R., Boulton, A. J. M., Feldman, E. L., Bril, V., Freeman, R., Malik, R. A., Sosenko, J. M., & Ziegler, D. (2017). Diabetic neuropathy: A position statement by the American Diabetes Association. *Diabetes Care*, *40*(1), 136–154. https://doi.org/10.2337/dc16-2042

Quan, H., Sundararajan, V., Halfon, P., Fong, A., Burnand, B., Luthi, J.-C., Saunders, L. D., Beck, C. A., Feasby, T. E., & Ghali, W. A. (2005). Coding algorithms for defining comorbidities in ICD-9-CM and ICD-10 administrative data. *Medical Care*, *43*(11), 1130–1139. https://doi.org/10.1097/01.mlr.0000182534.19832.83

Steyerberg, E. W., Vickers, A. J., Cook, N. R., Gerds, T., Gonen, M., Obuchowski, N., Pencina, M. J., & Kattan, M. W. (2010). Assessing the performance of prediction models: A framework for traditional and novel measures. *Epidemiology*, *21*(1), 128–138. https://doi.org/10.1097/EDE.0b013e3181c30fb2

Sun, X., & Xu, W. (2014). Fast implementation of DeLong's algorithm for comparing the areas under correlated receiver operating characteristic curves. *IEEE Signal Processing Letters*, *21*(11), 1389–1393. https://doi.org/10.1109/LSP.2014.2337313

Tani, M., Morimatsu, H., Takatsu, F., & Morita, K. (2012). The incidence and prognostic value of hypochloremia in critically ill patients. *The Scientific World Journal*, *2012*, 474185. https://doi.org/10.1100/2012/474185

Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B (Methodological)*, *58*(1), 267–288. https://doi.org/10.1111/j.2517-6161.1996.tb02080.x

Van Calster, B., McLernon, D. J., van Smeden, M., Wynants, L., & Steyerberg, E. W. (2019). Calibration: The Achilles heel of predictive analytics. *BMC Medicine*, *17*, 230. https://doi.org/10.1186/s12916-019-1466-7

van Houwelingen, H. C. (2007). Dynamic prediction by landmarking in event history analysis. *Scandinavian Journal of Statistics*, *34*(1), 70–85. https://doi.org/10.1111/j.1467-9469.2006.00529.x

Varma, S., & Simon, R. (2006). Bias in error estimation when using cross-validation for model selection. *BMC Bioinformatics*, *7*, 91. https://doi.org/10.1186/1471-2105-7-91

Vickers, A. J., & Elkin, E. B. (2006). Decision curve analysis: A novel method for evaluating prediction models. *Medical Decision Making*, *26*(6), 565–574. https://doi.org/10.1177/0272989X06295361

Youden, W. J. (1950). Index for rating diagnostic tests. *Cancer*, *3*(1), 32–35. https://doi.org/10.1002/1097-0142(1950)3:1<32::AID-CNCR2820030106>3.0.CO;2-3

Yuan, G.-X., Ho, C.-H., & Lin, C.-J. (2012). An improved GLMNET for L1-regularized logistic regression. *Journal of Machine Learning Research*, *13*, 1999–2030.

6 Supplementary material

- **S1. Four cohort definitions.** Four readings of "first ICU admission at first hospitalization", with counts and one-year mortality.
- **S2. Reference cohort: complete results.** Nested cross-validation of all nine configurations, LASSO selection frequencies, the model-selection steps, the secondary held-out estimate, thresholds, decision curves, coefficients and SHAP.
- **S3. Full baseline tables** of the reference and landmark cohorts.
- **S4. Comparison with the benchmark's Table 1** and the mixed-unit diagnostic for the white-cell differential.
- **S5. Hypertension definitions** and their associations with one-year death.
- **S6. LASSO cross-validation curve and coefficient path.**
- **S7. Exploratory single-split results** (all models, paired DeLong comparisons, thresholds, calibration, LASSO-only nesting).
- **S8. Comorbidity-timing recode**: acute codes removed and changes in each comorbidity.
- **S9. Implementation details.** Global random seed 42 and repeat seeds; LASSO implementation (λ grid of 100 values from $\lambda_{\max} = \max_j |\mathbf{x}_j^\top (\mathbf{y} - \bar{y})|/n$ down to $10^{-4}\lambda_{\max}$; liblinear intercept feature set to 1,000 so that the intercept penalty is negligible; binomial deviance as the criterion, whereas the benchmark's text names mean squared error but its figure shows binomial deviance); the benchmark's hyperparameters and their mapping to scikit-learn and xgboost; all tuning grids.
- **S10. Benchmark-style versus within-fold preprocessing**: LASSO selection, validation AUROC and SHAP.
- **S11. Code availability and reproducibility**: pipeline scripts, seeds and the frozen-model specification. All code, aggregate results and this manuscript's sources are available at https://github.com/hongfangyu123/BMI5101-DN-Mortality-P16; patient-level MIMIC-IV data are not included and require credentialed PhysioNet access.
- **S12. Exploratory extensions: complete results.** All 14 models (LR, RF, XGBoost and stacking with static, trajectory, burden and embedding inputs), all 45 paired comparisons, pooled out-of-fold calibration, decision-curve comparisons and trajectory-feature selection frequencies. The extension experiments were run on a separate machine; rebuilding their input features from our own data reproduced the same patients, outcomes and missingness pattern, with a maximum absolute difference of 2.8×10^{-14}.
- **S13. Post-hoc diagnostics of the held-out AUROC.** Exploratory analyses that did not establish the cause of the difference between nested-CV and held-out AUROC. The frozen LR's apparent training AUROC was 0.817. Resampling the OOF predictions 10,000 times to the held-out composition (136 deaths, 376 survivors) gave a mean AUROC of 0.800 (SD 0.021; 2.5th percentile 0.755), with 1.4% of resamples at or below 0.750; this ignores refitting and selection. The largest standardised mean difference between training and held-out sets across the 22 selected variables was 0.11; similar marginal distributions do not exclude differences in the joint distribution.